OXFORD

Detecting and correcting misclassified sequences in the large-scale public databases

Hamid Bagheri [1],*, Andrew Severin [2] and Hridesh Rajan [1]

[1] Dept. of Computer Science, Iowa State University, Ames, 50011, USA and
[2] Genome Informatics Facility, Iowa State University, Ames, 50011, USA.

* To whom correspondence should be addressed.

Associate Editor: XXXXXXX

Abstract

Motivation: As the cost of sequencing decreases, the amount of data being deposited into public repositories is increasing rapidly. Public databases rely on the user to provide metadata for each submission that is prone to user error. Unfortunately, most public databases, such as non-redundant (NR), rely on user input and do not have methods for identifying errors in the provided metadata, leading to the potential for error propagation. Previous research on a small subset of the non-redundant (NR) database analyzed misclassification based on sequence similarity. To the best of our knowledge, the amount of misclassification in the entire database has not been quantified. We propose a heuristic method to detect potentially misclassified taxonomic assignments in the NR database. We applied a curation technique and quality control to find the most probable taxonomic assignment. Our method incorporates provenance and frequency of each annotation from manually and computationally created databases and clustering information at 95% similarity.

Results: We found more than 2 million potentially taxonomically misclassified proteins in the NR database. Using simulated data, we show a high precision of 97% and a recall of 87% for detecting taxonomically misclassified proteins. The proposed approach and findings could also be applied to other databases.

Availability: Source code, dataset, documentation, Jupyter notebooks, and Docker container are available at https://github.com/boalang/nr.

Contact: hbagheri@iastate.edu

Supplementary information: Supplementary data are available at *Bioinformatics* online.

1 Introduction

Researchers use BLAST on the non-redundant (NR) database (NCBI (2018b)) on a daily basis to identify the source and function of a protein/DNA sequence. The non-redundant (NR) database encompasses protein sequences from non-curated (low quality) and curated (high quality) databases. It contains non-redundant sequences from GenBank translations (i.e. GenPept) together with sequences from other databases (Refseq (Pruitt *et al.* (2006)), PDB (Berman *et al.* (2003)), SwissProt (Boeckmann *et al.* (2003)), PIR (Wu *et al.* (2003)) and PRF). NR removes 100% identical sequences and merges the annotations and sequence IDs.

We have identified three root causes for annotation errors in the public databases: user metadata submission, contamination error in the biological samples, and computational methods. NCBI relies on the accuracy of the metadata provided by researchers that are depositing sequencing data into the database. Data are deposited into NCBI into Biosamples and Bioprojects as raw data, genome assemblies, and transcriptomes. Biosamples contain metadata describing the data type, scope, organism, publication, authors, and attributes, which include cultivar, biomaterial provider, collection date, tissue, developmental stage, geographical location, coordinates, and additional notes. This metadata is then propagated to the sequences that are deposited. For example, if data for DNA sequences were deposited by a plant researcher studying soybeans obtained from a soybean roots, then all sequences tied to that metadata will be labeled with the organism name *Glycine max*. If the researcher had in fact been working on *Glycine soja* then this would result in a misassignment of all *Glycine max* sequences.

The second key challenge that all large databases have is the issue of contamination (Schnoes *et al.* (2009)). For example, if the aforementioned

hypothetical soybean research did not remove the soybean root nodules during sample processing, then the tissue sample could also contain DNA from Nitrogen fixating soil bacteria that infect nodules leading to contamination of the sequences and ultimately the sequence database. NCBI is aware of the potential for contamination in sequence databases and describes potential sources of contamination that include: DNA recombination techniques (vectors, adaptors, linkers and PCR primers, transposon, and insertion sequences) and sample impurities (organelle, DNA/RNA, multiple organisms). NCBI encourages the use of programs to try to reduce issues with contamination. Specifically, they recommend screening for contamination using VecScreen (vec (2019)) and BLAST for the sequences used in sequencing library preparation. More broadly, they recommend BLAST to screen out bacterial, yeast, and *Escherichia coli* sequences and BLASTing against the NR database to identify potential contaminating sequences. Unfortunately, despite efforts to reduce contamination, sequences still end up in the NR database that is incorrectly taxonomically classified. This can limit our ability to identify contamination of future sequence submissions, as BLASTing against the database could propagate these types of errors as the database grows in size (Schnoes *et al.* (2009)). The contamination problem is not unique to NCBI but can be found in all large databases. A large-scale study of complete and draft bacterial and archaea genomes in the NCBI RefSeq database revealed that 2250 genomes are contaminated by human sequences (Breitwieser *et al.* (2019)). Breitwieser *et al.* reported 3437 spurious protein entries that are currently present in NR and TrEMBL protein databases.

The third key challenge is that there are errors in the annotations due to the computational error in tools that are based on homology to existing sequences to predict the annotations (Schnoes *et al.* (2009)). These errors have caused annotation accuracy and database quality issues over the years. Annotation errors are not limited to contamination or computationally predicted one. For instance, there exists evidence that suggests some of the experimentally derived annotations may be incorrect (Schnoes *et al.* (2009)).

Therefore, it will be beneficial for researchers to utilize a quality control method to detect misclassified sequences and propose the most probable taxonomic assignment.

To address these well-known problems, there are two approaches in the literature: phylogenetic-based approach and functional approach. For the first approach, Kozlov *et al.* (Kozlov *et al.* (2016)) have proposed a phylogeny-aware method to detect and correct misclassified sequences in public databases. They utilized the Evolutionary Placement Algorithm (EPA) to identify mislabeled taxonomic annotation. Edgar (Edgar (2018)) has studied taxonomy annotation error in rRNA databases. They showed that the annotation error rate in SILVA and Greengenes databases is about 17%. They also used the phylogenetic-based approach.

In the second approach, it is a common technique for quality control and data cleaning to utilize domain knowledge in the form of ontologies (Chu *et al.* (2015)). Gene Ontology (GO) (Ashburner *et al.* (2000)) has been suggested to infer aspects of protein function based on sequence similarity (Holliday *et al.* (2017)). The MisPred Nagy and Patthy (2013) and FixPred (Nagy and Patthy (2014)) programs are used to address the identification and correction of misclassified sequences in the public databases. The FixPred and MisPred methods are based on the principle that an annotation is likely to be erroneous if its feature violates our knowledge about proteins (Nagy *et al.* (2008)). MisPred (Nagy and Patthy (2013)) is a tool developed to detect incomplete, abnormal, or mispredicted protein annotations. There is a web interface to check the protein sequence online. MisPred uses protein-coding genes and protein knowledge to detect erroneous annotations at the protein function level. For example, they have found for a subset of protein databases that violation of domain integrity accounts for the majority of mispredictions. Modha *et al.* have proposed a pipeline to pinpoint taxonomic error as well as

identifying novel viral species (Modha *et al.* (2018)). There is another web-server for exploratory analysis and quality control of proteome-wide sequence search (Medlar *et al.* (2018)) that requires a protein sequence in a FASTA format. European Bioinformatics Institute (EMBL-EBI) developed InterPro (InterPro (2019)) to classify protein sequences at the superfamily, family and subfamily levels. UniProt has also developed two prediction systems, UniRule and the Statistical Automatic Annotation System (SAAS) (SAA (2019)), to annotate UniProtKB/TrEMBL protein database automatically. CDD is a Conserved Domain Database for the functional annotation of proteins (Marchler-Bauer *et al.* (2010)).

Exploring public sequence databases and curating annotations at large-scale is challenging. Previous research on the NR database focused on a small subset of the NR database and analyzed annotation error due to the computational requirements. There has been a study (Schnoes *et al.* (2009)) on misclassification levels for molecular function for a model set of 37 enzyme families. To the best of our knowledge, the amount of misclassification in the entire database has not been well quantified.

Here, we attempt to address these limitations in detecting and correcting annotations at large-scale and make the following contributions:

- We utilize a genomics-specific language, BoaG, that uses the Hadoop cluster (Bagheri *et al.* (2019)), to explore annotations in the NR database that is not available in other works.
- We also present a heuristic-based method to detect misclassified taxonomic assignments in the NR database that is low-cost and easy to use. We automatically generate a phylogenetic tree from a list of taxonomic assignments and use the tree, along with frequency, the provenance (database of origin) of each taxonomic annotation, and clustering information from NR at 95% similarity to identify potential misclassification and propose the most probable taxonomic assignment.
- The technique proposed in this work could be generalized to apply to other public databases and different kinds of annotations like protein functions. In this work, we address the taxonomic annotation error in protein databases. We also tested our approach on the RNA dataset introduced in the literature.

We have identified "29,175,336" proteins in the NR database that have more than one distinct taxonomic assignments, among which "2,238,230" (7.6%) are potentially taxonomically misclassified. We also found that the total number of potential misclassifications in clusters at 95% similarity, above the genus level, is "3,689,089" out of 88M clusters, which are 4% of the total clusters. This percentage of misclassifications in NR has a significant impact due to the potential for error propagation in the downstream analysis (Mukherjee *et al.* (2015)).

The rest of the paper is organized as follows. In Section 2, we present methods and materials for dataset generation and our approach. In section 3, we discuss the results of taxonomically misclassified proteins within sequences and in NR 95%. We also present the correcting approach for detected sequences in Section 4, we conclude with suggestions for the future.

2 Materials and methods

In this section, we will describe the overview architecture of our detection and correction approach. Then, we describe the dataset generation and how we generate a phylogenetic tree from taxonomic assignments. Next, we discuss our detection algorithm to find misclassified sequences. Then, we describe our approach to propose taxonomic assignments for the sequences identified as misclassified. Finally, we will describe the sensitivity analysis on changing the different parameters to propose the taxonomic assignments.
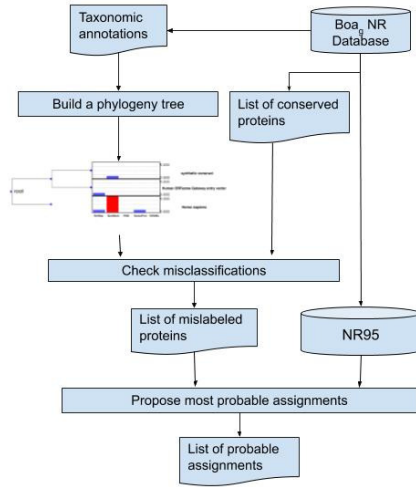
Fig. 1. Overview architecture of the proposed method to detecting taxonomically misclassified sequences in the NR database. Diagram shows the raw dataset and steps for the proposed work.

2.1 An overview of the method

Figure 1 shows an overview of our approach. The NCBI'NR database files were downloaded from NCBI (2018b) on Oct 20, 2018. Taxonomic information was obtained from XML files downloaded from the NCBI ftp server (NCBI (2018a)). CD-HIT (Fu *et al.* (2012)) (version v4.6.8-2017-1208) was used to cluster NR protein sequences into clusters at 95% similarity using the following parameters (-n 5 -g 1 -G 0 -aS 0.8 -d 0 -p 1 -T 28 -M 0). These parameters use a word length of 5 and require that the alignment of the short sequences is at least 80% of its length. The data acquisition, preprocessing, and clustering took about 3 days. The detection and correction part took about 8 hours.

We took the NR protein FASTA files that have the definition lines containing annotations from different databases and generate the BoaG format that took about 2 hours. Each definition line in the raw data includes protein ID, protein name followed by an organism name in square brackets, e.g., ">AAB18559 unnamed protein product [Escherichia coli str. K-12 substr. MG1655]". BoaG is a domain-specific language that uses a Hadoop based infrastructure for biological data (Bagheri *et al.* (2019)). A BoaG program is submitted to the BoaG infrastructure. It is compiled and executed on a distributed Hadoop cluster to execute a query on the BoaG-formatted database of the raw data. BoaG has aggregators that can be run on the entire database or a subset of the database taking advantage of protobuf-based schema design optimized for both the data and the computation. These aggregators are similar to but not limited to aggregators traditionally found in SQL databases and NoSQL databases like MongoDB. A BoaG script requires fewer lines of code, provides storage efficiency, and automatically parallelized large-scale analysis.

2.1.1 Dataset generation

To describe our dataset, let $\mathbb{D}$ denotes the protein and clustering dataset in our study: $\mathbb{D} = \{P, C, \tau, F\}$. Here, $P = \{P_1, P_2, \ldots, P_m\}$ is a set of all the proteins in the NR database. $C = \{C_1, C_2, \ldots, C_n\}$ represents a set of all clusters at 95% similarity. $|P|$ and $|C|$ in our dataset are about 174M and 88M respectively. τ is a set of taxonomic assignment for proteins, and F is a set of functions in the NR database. In this work, we focus on exploring taxonomic assignments.

Definition 1. *Cluster. We define cluster as a set of protein sequences such that their sequence are 95% similar and their sequence length is 80% similar. Every particular cluster, C_j, has k members:*

$$C_j = \{P_1, P_2, \ldots, P_k\}, and \, k \in [1, |P|] \tag{1}$$

In Definition 1, each protein sequence belongs to exactly one cluster at 95% similarity, and each cluster has one representative sequence. If a protein is not identical in sequence and length, it will fall into a cluster with no other member.

2.1.2 Generating phylogenetic tree from taxonomic assignments

We get the list of taxonomic assignments that originate from different databases (manually reviewed and computationally created) and build a phylogenetic tree by utilizing the ETE3 library (Huerta-Cepas *et al.* (2016)). This library utilizes the NCBI taxonomy database that is updated frequently.

Definition 2. *Annotation List. Each phylogeny tree is associated with one particular protein, P_i, and has the set of taxonomic assignments that originate from different databases. Here $A_{i,j}$ denotes annotation number j for protein P_i :*

$$\tau(P_i) = \{A_{i,1}, A_{i,2}, \ldots, A_{i,j}\}, j \in [1, |\tau|] \tag{2}$$

For example, the protein sequence AAB18559 has taxonomic assignments of "511145" and "723603" that each appeared once.

Definition 3. *Provenance. For the particular protein P_i, we define $prov(A_{i,a})$ the provenance of annotation $A_{i,a}$ as a set of databases that the annotation $A_{i,a}$ originates from:*

$$prov(A_{i,a}) \in \{GenBank, trEMBL, PDB, RefSeq, SwissProt\} \tag{3}$$

In Definition 3, annotations from GenBank, trEMBL, and PDB are calculated computationally, while annotations from RefSeq and SwissProt are manually reviewed. For example, $prov(511145) = GenBank$ meaning that the tax id "511145" for the sequence AAB18559 originates from the GenBank database.

Definition 4. *Annotation Probability. We define probability for each taxonomic assignment based on the frequency of each annotation divided by total taxonomic assignments from different databases as follows:*

$$prob(A_{i,a}) = \frac{freq\left(A_{i,a \in Comp}\right) + w \times freq\left(A_{i,a \in Rev}\right)}{\sum_{j \in Comp} freq\left(A_{i,j}\right) + \sum_{j \in Rev}(w \times freq\left(A_{i,j}\right))} \tag{4}$$

In Definition 4, $A_{i,a \in Comp}$ represents the annotation that calculated computationally (Comp) from databases i.e., GenBank, trEMBL, PDB, and $A_{i,a \in Rev}$ denotes the reviewed (Rev) one from RefSeq, SwissProt. One annotation might originate from both reviewed and computational created databases. We use a conservative weighting factor, w, to denote the importance of the experimental annotation (manually reviewed) in which w is an integer number and $w \geq 1$.

The upper bound for total proteins, i.e. $|P|$, is 174M at the time we downloaded NR. Each leaf node, V_a, in the phylogenetic tree is annotated with the information described in the Definitions 2, 3, and 4. There are list of frequencies and provenances, shows as top bar, since one particular taxonomic annotation might originate from different databases:

$$V_a = \{prob\left(A_{i,a}\right), \overline{freq\left(A_{i,a}\right), prov\left(A_{i,a}\right)}\}. \tag{5}$$

For particular protein P_i, we define most probable annotation (MPA) as $MPA(P_i) = A_{i,j}$ as an annotation with the highest probability among

the set of annotations. In addition, we define least probable annotation (LPA), with the lowest probability, that potentially might be misclassified as $LPA(P_i) = A_{i,k}$, in which $i \neq j$.

Definition 5. ***Conserved Proteins.*** *We define a conserved protein as a protein that has more than 10 distinct taxonomic assignment. List of these conserved proteins are shown in our repository (https://github.com/boalang/nr).*

$$P_i \; such \; that \; |\tau(P_i)| \geq 10 \tag{6}$$

2.2 Approach to detect taxonomic misclassification

Our approach is as follows: first, we run a BoaG query on the NR database. This query runs on the full NR database in the Hadoop cluster. BoaG query is shown in the supplemental file. The algorithm 1 describes the detection approach for misclassified sequences. It iterates over the entire NR database. In line 4, it takes a protein P_i and generates a phylogeny tree from the set of taxonomic assignments for P_i. Then, in line 5, it checks if it has a misclassification. If the lowest common ancestor (LCA) is the root level, it means there is a considerable distance between taxonomic assignments for that particular protein sequence. Therefore, there is a potential misassignment among the list of the taxonomic assignments due to the contamination in the sample, error in the computational method, or data entry by the researchers who deposited the sequence. We call this a root violation or conflict. We also consider superkingdom, phylum, class, order, and family violations. In addition, we looked at the highly conserved proteins to remove false positives because conserved proteins might appear in species that are far from each other, i.e., belong to different domains in the phylogeny tree. We did not remove the list of conserved proteins in the dataset, since they contain taxonomic information that were utilized for proposing taxonomic assignment for the misclassified sequences. Assume P_i belongs to C_j. Once we detected the violation in P_i, we look at the cluster C_j and consider the most frequent taxonomic assignment as the correct taxa. Details are shown in Section 2.3.

Algorithm 1 The NR misassignment detection algorithm. Input comes from the BoaG query (See supplemental)

1: **procedure** DetectMisassignments($\mathbb{D}$)
2:    $NRLength \leftarrow |P|$ ▷ m = 174M proteins
3:    **while** $i \leq NRLength$ **do**
4:       $phylo \leftarrow PhyloTree(P_i)$.
5:       **if** *misassigned(phylo)* && not conserved(P_i) **then**
6:          print (misassignment found in P_i)
7: **procedure** PhyloTree(P_i)
8:    $ncbi \leftarrow ncbiTAXA()$ ▷ used to generate phylogeny tree
9:    $phyloTree \leftarrow ncbi.get\_topology(P_i)$ ▷ From taxa list
10:    **for** $A_{i,a}$ in $\tau(P_i)$ **do**
11:       $V_a \leftarrow prob(A_{i,a}), list(freq(A_{i,a}), prov(A_{i,a}))$
12:    **return** phyloTree.

The algorithm 1 requires $\mathcal{O}(|P| * |\tau|)$ time. Here, $|P|$ is the size of proteins in the NR database and $|\tau|$ is the upper bound of number of taxonomic assignments per proteins. In line 5, $misassigned(phylo)$ verifies if the LCA of the generated tree shows a root violation or any other violations. The $conserved(P_i)$ expression checks if the protein sequence is a conserved one (See Eq. 6). This requires $\mathcal{O}(1)$ time because this is a straight-forward fetch, and we have the pointer to the root of the tree to check the LCA. In line 5, to check that a protein is not in a conserved list, Definition 5, it requires a membership test and takes $\mathcal{O}(1)$ time. This conserved list is a precomputed list from our dataset that is shown in our

repository. We wrote a multi-threaded Python code, and the total run time for the algorithm was seven hours for the entire NR database on an iMac (Retina 5K, 27-inch, Late 2015) with core i7 and 32 GB RAM. For the second procedure, in line 11, the algorithm requires $\mathcal{O}(|\tau|)$ to calculate the probability of each leaf in the generated phylogenetic tree.

Algorithm 2 Annotation correction: The Most Probable Annotation (MPA) for the misclassified sequences. Input from the BoaG query (See supplemental)

1: **procedure** mostProbable(P_i, p, c)
2:    $top\_ann \leftarrow max(prob(\tau(P_i)))$ ▷ Most probable taxa
3:    **if** $prob(top\_ann) \geq p$ **then**
4:       **return** (top_ann).
5:    **else**
6:       $cluster \leftarrow C_j$ in which $P_i \in C_j$
7:       $top\_ann \leftarrow ClusterMostProbable(cluster, p, c)$.
8:    **return** top_ann.
9: **procedure** ClusterMostProbable($clustr, p, c$)
10:    **if** $size(cluster) \geq c$ **then**
11:       **for** $A_{i,a}$ in $\tau(cluster)$ **do**
12:          $V_a \leftarrow prob(A_{i,a}), list(freq(A_{i,a}), prov(A_{i,a}))$
13:       $top\_ann \leftarrow max(prob(\tau(cluster)))$ ▷ Most probable taxa
14:       **if** $prob(top\_ann) \geq p$ **then**
15:          **return** top_ann
16:    **else**
17:       **return** $null$ ▷ Cannot fix misclassification

2.3 The most probable taxonomic assignment for detected misclassifications

For the detected misclassified sequences, we defined criteria to propose the most probable taxonomic assignment (MPA). First, we ran a BoaG query to retrieve the annotations and clustering information at 95% similarity. The BoaG script is shown in the supplemental file. As shown in Definition 4, we considered provenance or database of origin, frequency of annotations to calculate the probable taxonomic assignment (MPA), which is the highest probability. Let's assume that P_i belongs to cluster C_j. If the algorithm does not find the MPA within a certain threshold, probability p, then we look at the cluster of 95% similarity that the sequence belongs to. Second, we found the most probable taxonomic assignment in C_j. If a particular taxonomic assignment was the most frequent one in C_j then we return that annotation as the MPA for the protein sequence P_i. For example, in cluster C_j, seven sequences out of 10 sequences have a specific annotation. Then, we consider this annotation to be the most probable annotation protein sequence P_i with 70% confidence.

Details are shown in the algorithm 2. In line 2, for a particular protein P_i, it returns the most frequent taxonomic assignment within a certain threshold p. Let's assume we want a taxonomic assignment that appears more than 70% of the time. If the algorithm does not find the MPA, it checks the cluster C_j with 95% similarity that this sequence belongs to and finds the one with a certain probability, p, and a cluster size, c (line 7). In line 9, $ClusterMostProbable$ takes the cluster id and finds the most probable taxonomic assignment in the cluster (line 13).

The algorithm 2 requires $\mathcal{O}(|\tau(P)|)$ time, Definition 2, to find the top(1) or maximum probability of an annotation in the list of annotations.

2.4 Simulated and literature dataset

To evaluate the performance of our taxonomic misclassification approach, we generated a simulated dataset. We took a subset of one million proteins
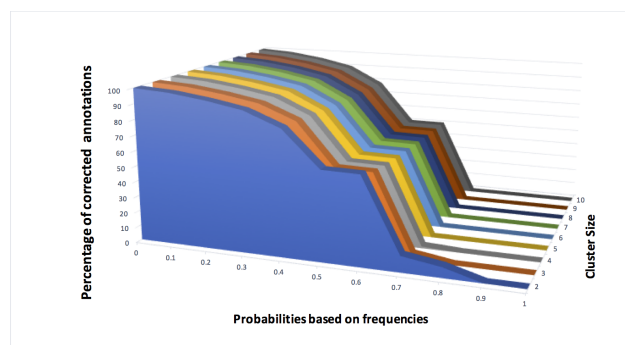
Fig. 2. The input parameters to the sensitivity analysis are probability based on the annotation frequencies and the cluster size. The Z-axis shows for what percentages of the misclassified sequences, our approach can propose taxonomic assignment from the clustering information.

Table 1. Detected misclassified taxonomic proteins in the NR database.

| taxa	total	root	Kingdom	Phylum	Class	Order	Family
2	**17,496,167**	**30,237**	**47,271**	**202,205**	59,606	177,132	290,065
3	**5,921,066**	**14,376**	**19,666**	**107,705**	38,575	104,709	236,515
4	2,132,971	4,673	21,587	64,801	17,662	47,914	94,054
5	1,022,482	3,143	9,469	34,322	10,050	27,295	53,276
6	642,760	2,509	5,662	24,136	7,333	23,324	37,998
7	388,794	1,572	3,959	12,972	5,905	13,488	27,221
8	262,682	1,121	2,803	5,988	5,375	10,075	16,340
9	190,756	783	2,647	3,825	3,173	7,557	12,681
10	156,767	667	1,843	3,805	2,451	6,413	11,327
>10	960,891	10,940	23,232	30,048	38,679	46,391	107,679

Fig. 3. Phylogenetic tree generated for sequence ID NP_001026909. Taxonomic assignments originate from GenBank, trEMBL, PDB, RefSeq, and SwissProt database.

of the reviewed dataset, i.e. RefSeq database, and randomly misclassified 50% of the proteins in the sample by adding a taxonomic assignment from another phylum or kingdoms. Then, we tested if the approach can detect these sequences. We also tested our approach for detecting misclassified sequences and correcting them on the real-world data, presented in the literature (Kozlov *et al.* (2016); Edgar (2018)). These works have focused on the RNA dataset, and they quantified misclassified RNA sequences. We also used CD-HIT to cluster RNA databases based on 95% sequence similarity. Further details on the simulated dataset, scripts, and data files can be accessed from https://github.com/boalang/nr.

2.5 Sensitivity analysis

We define sensitivity analysis as a way that an input parameter affects the output of the proposed approach. Here, probability based on annotation frequencies and the cluster size are the two input parameters that affect what percentages of detected misclassified sequences that we can fix, i.e., MPA, as shown in Algorithm 2 on the NR dataset. The algorithm will not give the same suggestion for changes in parameters. For example, if we change the cluster size, no. of proteins in the cluster, it may or may not find correct taxa. We conducted a sensitivity analysis based on the probability of each annotation that we defined in Definition 4 and the size of the cluster of 95% that the sequence belongs to. We run the algorithm to find the most probable taxonomic assignments (MPA) with different clusters size, c, and with different probabilities, p. As it is shown in figure 2, with a probability of 0.4 and without giving more weight to the annotations that verified experimentally, we could provide a most probable taxonomic assignment to about 60% of the proteins that we detected as misclassified. We also extended sensitivity analysis by giving more weight to the experimental taxonomic assignment with the probability of 0.4 we could provide the most probable taxonomic assignment for more than 80% of the sequences that were identified as a misclassification.

3 Results

In this section, we present the number of proteins that are misclassified taxonomically. We also present the performance of our work on the simulated dataset and the datasets presented in the literature. Then, we describe our findings on misassignments in the clusters. Next, we present correcting taxonomic misclassification. Finally, we discuss a case study that we explored deeply to identify a subset of clusters that contain sequences with a taxonomic misclassification.

3.1 Detected taxonomically misclassified proteins

We found "29,175,336" proteins in the NR database that have more than one distinct taxonomic assignments. The rest of the proteins have identical taxonomic assignments, even though they originate from different databases. The total number of potential taxonomically misclassified sequences is "2,238,230" out of "29,175,336" (7.6%) at the time of download. This percentage of NR is significant because of the error propagation in the downstream analysis (Mukherjee *et al.* (2015)). Table 1 shows the number of violations in the protein sequences in NR at the superkingdom to the family level that have been detected by applying distance in the phylogenetic tree. The second column shows the number of total proteins that have a certain number of taxonomic assignments. For example, there are "17,496,167" protein sequences in NR that have two taxonomic assignments in which "30,237" of them have potential root violations and "47,271", "202,205", "59,606", "177,132", "290,065" have kingdom, phylum, class, order, and family violations respectively. For the NR datasets, we did a sample study of 1000 samples, and manually found 5.5% misassignment. The potentially misclassified sequences detected by the approach was around 7.6% that is consistent with the total number that was manually found, i.e. 5.5%.

Table 1 shows proteins that have less than 10 taxonomic assignments. The last row shows all other proteins with more than 10 assignments. The first two bold rows show the highest potential misassignments because if a protein has two or three taxonomic assignments and shows a root or kingdom violation, it is more likely to be misclassified. The full list of detected misclassified proteins, and detailed analysis are shown in our GitHub repository. We did not report the genus conflict since the probability of a false-positive misclassification is much higher compared to other taxonomic levels of conflict, such as root and superkingdom.

Figure 3 shows one example of a detected misclassified protein, with an id of NP_001026909. Since the lowest common ancestor in this tree is the root, it means those taxonomic assignments belong to a different kingdom. Leaves are annotated with a frequency of each taxonomic assignment as a bar chart from all reviewed and unreviewed databases i.e., RefSeq (Pruitt *et al.* (2006)), GenBank (Benson *et al.* (2008)), PDB (Berman *et al.* (2003)), UniProt\SwissProt (Boeckmann *et al.* (2003)),

and UniProt\TrEMBL (Consortium (2014)) respectively. As it is shown in the annotations, there are potential misassignments even though the key IDs originate from the reviewed databases, i.e., RefSeq and SwissProt. In this example, *synthetic construct* is the misassignment, and the MPA for this protein is *homo sapiens*.

We also explored some clusters in depth as a case study and identified proteins that are taxonomically misclassified as *Glycine*, which are in fact contamination in the sample (see supplementary materials).

3.2 Performance on simulated and real-world dataset

Our approach to detecting taxonomically misclassified proteins on the simulated dataset showed 87% recall and 97% precision. We define *true positive* (TP) as sequences that misclassified in the sample, and our approach retrieves those sequences. *False positives* (FP) are sequences that do not have misassignments, but our approach classified them as misclassified sequences. *False negative* (FN) is a reviewed sequence which the algorithm incorrectly classifies as correct (not misclassified), while it is misclassified. Some of these false negatives are due to changes in the taxonomies over time. Some taxonomic IDs might be obsolete, deleted, or get merged into other tax ids. We also found that some of the trees generated by NCBI API have the root named "Cellular Organisms" with rank equal to "no rank", that did not fall in any of the taxonomic ranking. We use the following formula to calculate precision and recall:

$$precision = \frac{TP}{TP + FP}; recall = \frac{TP}{TP + FN} \qquad (7)$$

We extended our experiment and added more than two random assignments to the proteins and the precision increased. The reason is that adding more random assignments increases the distance among tax IDs in the phylogeny tree and hence increases the chance of detection by the approach. We also tested our approach on the dataset presented by (Edgar (2018)) in which they explored the Greengenes and the SILVA database for taxonomic error. Our methods reproduced their finding on annotation conflicts among SILVA and Greengenes (McDonald *et al.* (2012)) database. We did not run their approach on the simulated dataset since it was designed to detect misassignments in rRNA sequences, not proteins. For evaluating our work, we looked for similar works that focused on detecting taxonomic misassignments. However, their approach was hard-coded for RNA sequences. Therefore, we modified our approach to test on their dataset. The proposed work focuses on inconsistencies among the list of taxonomies, and it can be applied to the RNA sequences as well. We clustered their dataset at 95% similarity and used the same consensus-based technique to detect conflicts between sequences and clusters. The phylogeny-aware technique proposed by Kozlov *et al.* , called SATIVA, identifies and corrects misclassified sequences for RNA databases (Kozlov *et al.* (2016)). They utilized the Evolutionary Placement Algorithm (EPA) to detect misclassified sequences. In their approach, a reference tree is created. Then, to estimate the most likely placements of the query sequence in the reference tree, they use EPA. We took their RNA dataset and cluster the sequences at 95% similarity, then utilized our technique to check if the annotation of each sequence has a conflict with a cluster that the sequence belongs to. There is a difference between the NR dataset and the RNA dataset presented by Kozlov *et al.* in terms of the number of taxonomic annotation. In their experiment, they have one taxonomic label for each sequence; however, in the NR database, there are several annotations for each protein sequence. Therefore, their technique is not designed to detect misclassification in a set of given annotations. In terms of running time, the clustering at 95% is less expensive than running sequence alignment and generating phylogeny-tree and verifying each query sequence. Therefore, our approach is scalable for large-scale sequence databases. In general,

Table 2. Accuracy of detecting misassignments and the comparison with work presented in SATIVA (Kozlov et al. (2016))

Precision		Recall		Runtime	
SATIVA	Proposed	SATIVA	Proposed	SATIVA	Proposed
0.93	**0.98**	**0.98**	0.90	116 min	**12 min**

The **best** values are highlighted.

Table 3. Misclassification in clusters of the NR database at 95% similarity.

| # taxa	total	Root	kingdom	Phylum	Class	Order	Family
2	12,960,476	17,099	92,526	263,844	100,560	267,251	461,795
3	4,683,663	9,825	39,940	153,678	63,414	153,996	291,418
4	2,328,246	7,314	25,361	95,038	33,603	102,671	173,810
5	1,293,767	5,136	17,915	56,510	22,253	62,025	101,675
6	566,574	4,936	14,660	39,410	15,738	46,913	66,741
7	566,574	3,652	13,642	23,206	12,160	40,760	49,046
8	403,513	2,719	8,433	10,622	10,577	24,259	36,463
9	289,289	1,635	6,655	7,291	8,890	19,608	28,549
10	235,451	1,423	4,744	8,991	8,586	16,070	22,026
>10	1,832,313	22,921	63,513	3,951	65,196	155,804	200,642

Fig. 4. Misclassification detected in the cluster ID 490503 at 95% similarity of the NR dataset.

examining the distance on the phylogenetic tree of multiple annotations for the shorter sequences performs better compared to the alignment-based approaches with the reference databases. Table 2 shows the standard values for precision and recall, as well as the running time comparison. Our approach to detect misassignments on the sample RNA dataset has a lower recall. This is due to the relatively smaller datasets that caused some clusters to have few members and made it challenging to detect misclassified sequences.

3.3 Detected misassignments in clusters

Table 3 shows the number of potential misclassified sequences in the clusters at 95% similarity at the superkingdom level to the family level that has been detected by applying distance in the phylogenetic tree. The second column shows the number of entire clusters that have a certain number of taxonomic assignments. For example, there are "12,960,476" clusters at 95% similarity that have two taxonomic assignments in which "17,099" of them have potential root violations and "92,526", "263,844", "100,560", "267,251", "461,795" have kingdom, phylum, class, order, and family violations respectively. The number of root violations for two tax assignments in clusters is less than sequences because there are protein sequences that do not belong to any clusters at 95% similarity. 64M out of 174M proteins (36%) in the NR database are unclustered (see supplemental files). The total number of potential misclassifications for clusters at 95% similarity, without genus level, is "3,689,089" out of "25,159,866" clusters that have more than one taxa, which are 15% of total clusters.

Figure 4 shows an example of detected taxonomically assigned annotations in the cluster-id 490503. Leaves are annotated with databases of origins and frequency of taxonomic assignments as a bar chart from all reviewed and unreviewed databases, i.e., RefSeq, GenBank, PDB UniProt\SwissProt, and UniProt\TrEMBL respectively. For this cluster, there is no annotation from PDB, SwissProt, and TrEMBL databases.

Fig. 5. Conflicts between the set of top three taxonomic assignment in sequences and the top three taxonomic assignment in the clusters of the NR database.

Table 4. Proposed taxa for the detected misclassified sequences in NR. Last column shows the confidence score (CS).

Protein ID	Cluster ID	Original taxa	Proposed taxa	CS
AAB18559	18982245	uncultured actinobacterium	Escherichia coli	1
AAT83007	21005513	Mycobacteroides abscessus	Cutibacterium acnes	0.8
CCW09133	9901357	Streptococcus pneumoniae	Bacillus cereus	0.5
KFV03115	13041247	Tauraco erythrolophus	Pelodiscus sinensis	0.4
YP_950729	83178931	Staphylococcus virus PH15	firmicutes	0.8

3.4 Correcting Taxonomic Misclassification

Each protein sequence belongs to one and only one cluster. We analyzed the set of top three taxonomic annotations of each sequence and compared them with the top three taxonomic annotations of the cluster the sequence belongs to. For example, top three taxonomic assignment for sequence with id AAA32344 is '10743', '1182665', '656390'. This sequence falls in the cluster-id 8461728, and the top tax ids in this cluster are '562', '83334', '621'. We consider this as a conflict between sequence AAA32344 and cluster 8461728. All three annotations are different; therefore, we consider this case as three conflicts. If two annotations out of three are different, we classify this as *two conflicts*. If one taxonomic annotation is different from the two sets, we classify it as *one conflict*. Finally, if the three annotations are identical, there is no conflict. Figure 5 shows different percentages of conflicts from the subset of one million sequences in the NR database.

Table 4 shows several examples of the protein sequences that we have found to be misclassified in the NR database. The first column represents the sequence id, and the second column is the cluster id corresponds to the sequence. The third column shows the original taxonomic assignment, and the forth column is the proposed taxonomy based on the consensus information from the clusters of the NR database at 95% similarity. The last column is Confidence Score (CS), a number between 0 and 1, shows how confident we are in proposing new taxonomic assignment based on the consensus information from the clusters at 95% similarity. This score calculated from the clusters' information as top taxonomic assignment, i.e. most frequent one, in the cluster divided by total taxa in the cluster. The assumption here is that the consensus of multiple independent sequence annotations can catch simple misannotation errors. For example, protein sequence with id YP_950729 has *Staphylococcus virus PH15* as its taxonomic assignment. It falls in cluster id 83178931 and the recommended annotation is *firmicutes*. We also conducted similar analysis on the dataset by SATIVA, and could reproduce the proposed taxa based on the consensus information from the clusters. For the dataset by Edgar (Edgar (2018)) since the number of sequences was small, we could not get clusters with enough members to suggest annotations.
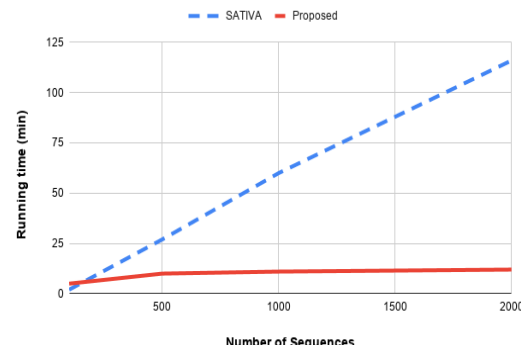

Fig. 6. Compare running time of the proposed work with the SATIVA method. We used dataset from the SATIVA paper.

3.5 Running time

We conducted an analysis of the RNA dataset presented by SATIVA (Kozlov *et al.* (2016)) with different samples of sequences. Firstly, we took 100 sequences and ran SATIVA in the sample. Next, we took 500 sequences. In two other experiments, we took 1000 and 2000 additional sequences and recorded the running time. Figure 6 shows the comparison in terms of running time (min) between proposed work and the SATIVA method. The most time-consuming part of our approach is the clustering time (run by CD-HIT software). By adding more sequences, the runtime slightly increased. In contrast, for the SATIVA method, as we increase the number of sequences, the running time increases significantly. The computational expensive part of the SATIVA approach is the phylogenetic methods (the Evolutionary Placement Algorithm) it employs. The comparison between the proposed approach and SATIVA method has been made on the local system iMac (Retina 5K, 27-inch, Late 2015) with core i7 and 32 GB RAM.

4 Discussion and conclusion

In this work, we addressed taxonomically misclassified sequences in the large publicly available databases by utilizing our domain-specific language and Hadoop-based infrastructure. We focused on the misassignments at the taxonomic level, and similar to MisPred (Nagy and Patthy (2013)), we utilize the current knowledge of organismal classification, to detect annotation errors. Similar to (Holliday *et al.* (2017)), we utilized this form of knowledge-based reasoning for quality control and detect annotation errors.

Compared to other works, our work differs in that we do not need to run sequence similarity to explore annotations and find taxonomic inconsistency for each query sequence in the NR database. Instead, first, we clustered the NR proteins at the data generation phase, and this is a one-time cost and used the clustering information later to detect annotation error and propose the most probable annotations. In this work, we proposed a heuristic method to find inconsistencies in the metadata, i.e., taxonomic assignments. In our method, we proposed the most probable taxonomic assignment for each protein sequence. We applied this method to the entire database. We also provided a Python implementation in a that could be used for analyzing a list of annotations for any protein of interest and find the misclassification. The violations reported in this paper in Table 1 are the upper bound of the misassignments. The more stringent filter includes *hypothetical protein* and *membrane protein* functions in the list of conserved protein, which will lower the number of identified misclassification.

We use open-source CD-HIT clustering software only at the data generation phase, and we could utilize any other clustering software. Steinegger *et al.* have built a novel clustering tool that clusters a huge protein database in linear time(Steinegger and Söding (2018)). Since this one-time clustering cost happens only in the data generation phase, our approach to detect misassignments and propose the most probable taxonomic assignment is scalable.

4.1 Applications and limitations

At 95% similarity, 64M sequences in the NR remain unclustered. Therefore, if a particular protein remains unclustered, there is not enough consensus information to correct annotation for that protein. A solution for this might be to take the Evolutionary Placement Algorithm (EPA) approach (Kozlov *et al.* (2016)) for these sequences that remains as future work. The proposed technique to detect misassignments may fail with recent horizontal gene transfer (HGT) events since HGT is not transferred from parent to offspring. However, the consensus information from the clusters might reveal annotation errors. The proposed heuristic technique and findings could also be applied to other databases. Current work focuses on detecting and correcting misassignments at the level of taxonomic assignments, and we do not address protein function annotations.

4.2 Conclusion

Misclassification can lead to significant error propagation in the downstream analysis. In this work, we proposed a heuristic approach to detect misclassified taxonomic assignments and find the most probable annotations for misclassified sequences. This method will be a valuable tool in cleaning up on large-scale public databases. The technique we proposed could be extended in the form of ontologies to address other annotation errors like protein functions.

5 Acknowledgments

References

(2019). Interpro: Classification of protein families. http://www.ebi.ac.uk/interpro/. Accessed: 2019-06-10.

(2019). Saas (statistical automatic annotation system). https://www.uniprot.org/help/saas. Accessed: 2019-06-10.

(2019). Vecscreen. https://www.ncbi.nlm.nih.gov/tools/vecscreen/. Accessed: 2019-06-10.

Ashburner, M., Ball, C. A., Blake, J. A., Botstein, D., Butler, H., Cherry, J. M., Davis, A. P., Dolinski, K., Dwight, S. S., Eppig, J. T., *et al.* (2000). Gene ontology: tool for the unification of biology. *Nature genetics*, **25**(1), 25.

Bagheri, H., Muppirala, U., Masonbrink, R., Severin, A. J., and Rajan, H. (2019). Shared data science infrastructure for genomics data, doi: https://doi.org/10.21203/rs.2.4295/v3. *BMC Bioinformatics*.

Benson, D. A., Karsch-Mizrachi, I., Lipman, D. J., Ostell, J., and Sayers, E. W. (2008). Genbank. *Nucleic acids research*, **37**(suppl_1), D26–D31.

Berman, H. M., Bourne, P. E., Westbrook, J., and Zardecki, C. (2003). The protein data bank. In *Protein Structure*, pages 394–410. CRC Press.

Boeckmann, B., Bairoch, A., Apweiler, R., Blatter, M.-C., Estreicher, A., Gasteiger, E., Martin, M. J., Michoud, K., O'donovan, C., Phan, I., *et al.* (2003). The swiss-prot protein knowledgebase and its supplement trembl in 2003. *Nucleic acids research*, **31**(1), 365–370.

Breitwieser, F. P., Pertea, M., Zimin, A. V., and Salzberg, S. L. (2019). Human contamination in bacterial genomes has created thousands of spurious proteins. *Genome research*, **29**(6), 954–960.

Chu, X., Morcos, J., Ilyas, I. F., Ouzzani, M., Papotti, P., Tang, N., and Ye, Y. (2015). Katara: A data cleaning system powered by knowledge bases and crowdsourcing. In *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data*, pages 1247–1261.

Consortium, U. (2014). Uniprot: a hub for protein information. *Nucleic acids research*, **43**(D1), D204–D212.

Edgar, R. (2018). Taxonomy annotation and guide tree errors in 16s rrna databases. *PeerJ*, **6**, e5030.

Fu, L., Niu, B., Zhu, Z., Wu, S., and Li, W. (2012). Cd-hit: accelerated for clustering the next-generation sequencing data. *Bioinformatics*, **28**(23), 3150–3152.

Holliday, G. L., Davidson, R., Akiva, E., and Babbitt, P. C. (2017). Evaluating functional annotations of enzymes using the gene ontology. In *The Gene Ontology Handbook*, pages 111–132. Humana Press, New York, NY.

Huerta-Cepas, J., Serra, F., and Bork, P. (2016). Ete 3: reconstruction, analysis, and visualization of phylogenomic data. *Molecular biology and evolution*, **33**(6), 1635–1638.

Kozlov, A. M., Zhang, J., Yilmaz, P., Glöckner, F. O., and Stamatakis, A. (2016). Phylogeny-aware identification and correction of taxonomically mislabeled sequences. *Nucleic Acids Research*, **44**(11), 5022–5033.

Marchler-Bauer, A., Lu, S., Anderson, J. B., Chitsaz, F., Derbyshire, M. K., DeWeese-Scott, C., Fong, J. H., Geer, L. Y., Geer, R. C., Gonzales, N. R., *et al.* (2010). Cdd: a conserved domain database for the functional annotation of proteins. *Nucleic acids research*, **39**(suppl_1), D225–D229.

McDonald, D., Price, M. N., Goodrich, J., Nawrocki, E. P., DeSantis, T. Z., Probst, A., Andersen, G. L., Knight, R., and Hugenholtz, P. (2012). An improved greengenes taxonomy with explicit ranks for ecological and evolutionary analyses of bacteria and archaea. *The ISME journal*, **6**(3), 610.

Medlar, A. J., Törönen, P., and Holm, L. (2018). Aai-profiler: fast proteome-wide exploratory analysis reveals taxonomic identity, misclassification and contamination. *Nucleic acids research*, **46**(W1), W479–W485.

Modha, S., Thanki, A. S., Cotmore, S. F., Davison, A. J., and Hughes, J. (2018). Victree: an automated framework for taxonomic classification from protein sequences. *Bioinformatics*, **34**(13), 2195–2200.

Mukherjee, S., Huntemann, M., Ivanova, N., Kyrpides, N. C., and Pati, A. (2015). Large-scale contamination of microbial isolate genomes by illumina phix control. *Standards in genomic sciences*, **10**(1), 18.

Nagy, A. and Patthy, L. (2013). Mispred: a resource for identification of erroneous protein sequences in public databases. *Database*, **2013**.

Nagy, A. and Patthy, L. (2014). FixPred: a resource for correction of erroneous protein sequences. *Database*, **2014**.

Nagy, A., Hegyi, H., Farkas, K., Tordai, H., Kozma, E., Bányai, L., and Patthy, L. (2008). Identification and correction of abnormal, incomplete and mispredicted proteins in public databases. *BMC Bioinformatics*, **9**(1), 353.

NCBI (2018a). NCBI XML for NR. https://ftp.ncbi.nlm.nih.gov/blast/temp/DB_XML/. Accessed: Oct 2018.

NCBI (2018b). Non-redundant database (nr). ftp://ftp.ncbi.nlm.nih.gov/blast/db/FASTA/. Accessed: Oct 2018.

Pruitt, K. D., Tatusova, T., and Maglott, D. R. (2006). Ncbi reference sequences (refseq): a curated non-redundant sequence database of genomes, transcripts and proteins. *Nucleic acids research*, **35**(suppl_1), D61–D65.

Schnoes, A. M., Brown, S. D., Dodevski, I., and Babbitt, P. C. (2009). Annotation error in public databases: misannotation of molecular function in enzyme superfamilies. *PLoS computational biology*, **5**(12), e1000605.

Steinegger, M. and Söding, J. (2018). Clustering huge protein sequence sets in linear time. *Nature communications*, **9**(1), 1–8.

Wu, C. H., Yeh, L.-S. L., Huang, H., Arminski, L., Castro-Alvear, J., Chen, Y., Hu, Z., Kourtesis, P., Ledley, R. S., Suzek, B. E., *et al.* (2003). The protein information resource. *Nucleic acids research*, **31**(1), 345–347.